

```
set.seed(3938425)
```

```
install.packages("caret")
```

```
Installing package into '/usr/local/lib/R/site-library'  
(as 'lib' is unspecified)
```

```
also installing the dependencies 'listenv', 'parallelly', 'future', 'globals', 'shape', 'future.apply', 'numDeriv',
```



```
install.packages("randomForest")
```

```
Installing package into '/usr/local/lib/R/site-library'  
(as 'lib' is unspecified)
```

```
library(caret)
```

```
➡ Loading required package: lattice
```

```
library(mgcv)
```

```
qsar <- readRDS("qsar.Rda")
```

```
set.seed(42)
```

```
train_indices <- sample(1:nrow(qsar), size = 700)
```

```
test_indices <- which(!(1:nrow(qsar) %in% train_indices))
```

```
predictors_train <- qsar[train_indices, 1:41] # First 41 columns are predictors  
response_train <- qsar[train_indices, 42]
```


```
set.seed(123)
```

```
predictors_test <- qsar[test_indices, 1:41] # First 41 columns are predictors  
response_test <- qsar[test_indices, 42]
```

```
trained_model <- train(  
  x = predictors_train, # predictors_train: Training predictors  
  y = response_train,   # response_train: Training response  
  method = "gbm",       # Method: Gradient Boosting Machine  
  tuneGrid = grid,      # Parameter grid for tuning  
  distribution = "bernoulli", # Distribution for classification  
  trControl = trainControl(method = "cv", number = 10)  
)
```

240	0.6362	-nan	0.0100	0.0001
260	0.6176	-nan	0.0100	0.0003
280	0.6015	-nan	0.0100	0.0001
300	0.5849	-nan	0.0100	0.0000
320	0.5705	-nan	0.0100	0.0001
340	0.5578	-nan	0.0100	0.0002
360	0.5458	-nan	0.0100	0.0001
380	0.5348	-nan	0.0100	-0.0001
400	0.5239	-nan	0.0100	-0.0000
420	0.5142	-nan	0.0100	-0.0002
440	0.5042	-nan	0.0100	-0.0001
460	0.4950	-nan	0.0100	-0.0002
480	0.4857	-nan	0.0100	-0.0001
500	0.4771	-nan	0.0100	0.0001
520	0.4692	-nan	0.0100	-0.0000
540	0.4618	-nan	0.0100	0.0000
560	0.4545	-nan	0.0100	-0.0002
580	0.4468	-nan	0.0100	0.0000
600	0.4393	-nan	0.0100	-0.0000
620	0.4320	-nan	0.0100	0.0000
640	0.4258	-nan	0.0100	-0.0001
660	0.4196	-nan	0.0100	-0.0001
680	0.4138	-nan	0.0100	-0.0002
700	0.4075	-nan	0.0100	-0.0001
720	0.4020	-nan	0.0100	-0.0001
740	0.3967	-nan	0.0100	0.0000
760	0.3915	-nan	0.0100	-0.0001
780	0.3861	-nan	0.0100	-0.0000
800	0.3810	-nan	0.0100	-0.0001
820	0.3754	-nan	0.0100	-0.0001
840	0.3706	-nan	0.0100	-0.0000
860	0.3654	-nan	0.0100	-0.0001
880	0.3606	-nan	0.0100	-0.0001
900	0.3562	-nan	0.0100	-0.0002
920	0.3516	-nan	0.0100	-0.0001
940	0.3473	-nan	0.0100	-0.0001
960	0.3429	-nan	0.0100	-0.0001
980	0.3385	-nan	0.0100	-0.0002
1000	0.3346	-nan	0.0100	-0.0001

```
plot(trained_model)
```



Q) Describe the main effects of the shrinkage, n.trees and interaction.depth parameters on the accuracy of the model. Also describe their possible interactions

Ans:- Shrinkage: Shrinkage controls the learning rate of the boosting process. A smaller shrinkage means each tree contributes less to the final prediction. Generally, smaller shrinkage values tend to improve accuracy, but they also require more computational resources and training time. However, too small a shrinkage value can lead to overfitting, especially if n.trees is large.

n.trees: specifies the number of boosting iterations. Increasing the number of trees usually leads to better model performance.

Interaction depth: Interaction depth controls the depth of interaction between variables in the model. A higher interaction depth allows the model to capture more complex interactions between predictors, potentially leading to better performance. It's essential to carefully tune this parameter to find the optimal balance between model complexity and generalization performance.

Interactions:

Shrinkage and n.trees: These parameters often interact, as a smaller shrinkage value requires more trees to achieve the same level of accuracy. Thus, the optimal combination of shrinkage and n.trees depends on the specific dataset and the trade-off

between computational resources and model performance. Interaction depth and n.trees: Increasing the interaction depth may require more trees to capture the additional complexity introduced by interactions between predictors. Thus, the optimal

combination of interaction depth and n.trees also depends on the dataset and the desired level of model complexity. Shrinkage and interaction depth: Lower shrinkage values may allow for deeper interactions between variables, potentially influencing the optimal interaction depth. However, too low a shrinkage value can lead to overfitting, so it's essential to balance these parameters carefully.

Q) If you look at the effect of interaction.depth, what would you conclude about the possible presence of interactions in the QSAR dataset?

Ans:- Improvement with increasing interaction.depth: the model's performance (e.g., accuracy) improves as the interaction.depth increases, it suggests that the dataset contains complex interactions between predictors.

```
best_params <- trained_model$bestTune
```

```
library(gbm)
```

```
Loaded gbm 2.1.9
```

```
This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com/gbm-developers
```

```
response_train_binary <- ifelse(response_train == "RB", 1, 0)
```

```
gbm_model <- gbm(
  formula = response_train_binary ~ ., # Define the formula for the model
  data = predictors_train,             # Training data
  distribution = "bernoulli",          # Distribution for classification
  n.trees = best_params$n.trees,       # Optimal number of trees
  interaction.depth = best_params$interaction.depth, # Optimal interaction depth
  shrinkage = best_params$shrinkage    # Optimal shrinkage
)
```

```
print(gbm_model)
```

```
gbm(formula = response_train_binary ~ ., distribution = "bernoulli",
     data = predictors_train, n.trees = best_params$n.trees, interaction.depth = best_params$interaction.depth,
     shrinkage = best_params$shrinkage)
A gradient boosted model with bernoulli loss function.
1000 iterations were performed.
There were 41 predictors of which 35 had non-zero influence.
```

```
library(randomForest)
library(caret)

single_tree <- train(x = predictors_train, y = response_train, method = "rpart")

bagged <- train(x = predictors_train, y = response_train, method = "treebag")

random_forest <- train(x = predictors_train, y = response_train, method = "rf")

boosted_default <- train(x = predictors_train, y = response_train, method = "gbm")

pred_single_tree <- predict(single_tree, newdata = predictors_test, type = "prob")

pred_bagged <- predict(bagged, newdata = predictors_test, type = "prob")

pred_random_forest <- predict(random_forest, newdata = predictors_test, type = "prob")

pred_boosted_default <- predict(boosted_default, newdata = predictors_test, type = "prob")

brier_score_single_tree <- mean((response_test - pred_single_tree[, "RB"])^2)
misclassification_rate_single_tree <- mean(ifelse(response_test == "RB", 1, 0) != apply(pred_single_tree, 1, which.max))

brier_score_bagged <- mean((response_test - pred_bagged[, "RB"])^2)
misclassification_rate_bagged <- mean(ifelse(response_test == "RB", 1, 0) != apply(pred_bagged, 1, which.max))

brier_score_random_forest <- mean((response_test - pred_random_forest[, "RB"])^2)
misclassification_rate_random_forest <- mean(ifelse(response_test == "RB", 1, 0) != apply(pred_random_forest, 1, which.max))

brier_score_boosted_default <- mean((response_test - pred_boosted_default[, "RB"])^2)
misclassification_rate_boosted_default <- mean(ifelse(response_test == "RB", 1, 0) != apply(pred_boosted_default, 1, which.max))

cat("Brier Score for Single Tree:", brier_score_single_tree, "\n")
cat("Misclassification Rate for Single Tree:", misclassification_rate_single_tree, "\n")

cat("Brier Score for Bagged Ensemble:", brier_score_bagged, "\n")
cat("Misclassification Rate for Bagged Ensemble:", misclassification_rate_bagged, "\n")

cat("Brier Score for Random Forest Ensemble:", brier_score_random_forest, "\n")
cat("Misclassification Rate for Random Forest Ensemble:", misclassification_rate_random_forest, "\n")

cat("Brier Score for Boosted Ensemble (Default Settings):", brier_score_boosted_default, "\n")
cat("Misclassification Rate for Boosted Ensemble (Default Settings):", misclassification_rate_boosted_default, "\n")
```

60	0.3978	-nan	0.1000	0.0003
80	0.3336	-nan	0.1000	-0.0007
100	0.2869	-nan	0.1000	-0.0013
120	0.2488	-nan	0.1000	-0.0011
140	0.2162	-nan	0.1000	-0.0013
150	0.2000	-nan	0.1000	-0.0002

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	1.2314	-nan	0.1000	0.0280
2	1.1806	-nan	0.1000	0.0222
3	1.1364	-nan	0.1000	0.0200
4	1.0914	-nan	0.1000	0.0220
5	1.0579	-nan	0.1000	0.0136
6	1.0164	-nan	0.1000	0.0168
7	0.9782	-nan	0.1000	0.0174
8	0.9439	-nan	0.1000	0.0151
9	0.9148	-nan	0.1000	0.0088
10	0.8910	-nan	0.1000	0.0106
20	0.7271	-nan	0.1000	0.0044
40	0.5846	-nan	0.1000	-0.0018
60	0.5035	-nan	0.1000	-0.0012
80	0.4523	-nan	0.1000	-0.0020
100	0.4044	-nan	0.1000	-0.0019
120	0.3682	-nan	0.1000	-0.0009
140	0.3366	-nan	0.1000	-0.0011
150	0.3199	-nan	0.1000	-0.0003

```
Warning message in Ops.factor(response_test, pred_single_tree[, "RB"]):
“-” not meaningful for factors”
Warning message in Ops.factor(response_test, pred_bagged[, "RB"]):
“-” not meaningful for factors”
Warning message in Ops.factor(response_test, pred_random_forest[, "RB"]):
“-” not meaningful for factors”
Warning message in Ops.factor(response_test, pred_boosted_default[, "RB"]):
“-” not meaningful for factors”
Brier Score for Single Tree: NA
Misclassification Rate for Single Tree: 0.8873239
Brier Score for Bagged Ensemble: NA
Misclassification Rate for Bagged Ensemble: 0.9492958
Brier Score for Random Forest Ensemble: NA
Misclassification Rate for Random Forest Ensemble: 0.9352113
Brier Score for Boosted Ensemble (Default Settings): NA
Misclassification Rate for Boosted Ensemble (Default Settings): 0.9380282
```

```
predicted_probs <- predict(gbm_model, newdata = predictors_test, type = "response")
```

```
brier_score <- mean((response_test - predicted_probs)^2)
```

```
predicted_labels <- ifelse(predicted_probs >= 0.5, 1, 0)
```

```
misclassification_rate <- mean(predicted_labels != response_test)
```

```
cat("Brier Score:", brier_score, "\n")
cat("Misclassification Rate:", misclassification_rate, "\n")
```

```
Using 1000 trees...
```

```
Warning message in Ops.factor(response_test, predicted_probs):
“-” not meaningful for factors”
Brier Score: NA
Misclassification Rate: 1
```

